

# Speeding up static analysis with clang-tidy-cache

January 27, 2021

# Introduction

# clang-tidy

- “a clang-based C++ static analysis tool<sup>1</sup>”
  - an extensible framework for diagnosing and fixing typical programming errors,
  - has a comprehensive suite of built-in checks for:
    - style violations,
    - language misuse,
    - anti-patterns,
    - common bugs,
    - etc.
  - provides a convenient interface for writing new checks,
  - is configurable, with a large set of options.

---

<sup>1</sup><https://clang.llvm.org/extra/clang-tidy/>

# Using clang-tidy with cmake

- `cmake` has built-in support for `clang-tidy`:
  - the `CXX_CLANG_TIDY` target property<sup>2</sup>.
  - the generated build system code includes instructions to run `clang-tidy`, typically chained with the compilation commands.

<sup>2</sup><https://cmake.org/cmake/help/latest/manual/cmake-properties.7.html> □ ▸ ◀ ◻ ▸ ◀ ≡ ▸ ◀ ≡ ▸ ≡ 🔍 ↺

Find the clang-tidy command / executable path:

```
find_program(  
  CLANG_TIDY_COMMAND  
  clang-tidy  
)
```

Add executable;

```
add_executable(my_target my_target.cpp)
```

or library target:

```
add_library(my_target my_target.cpp)
```

If `clang-tidy` was found, tell `cmake` to check the sources as a part of compilation:

```
if(CLANG_TIDY_COMMAND)
  set_target_properties(
    my_target PROPERTIES
    CXX_CLANG_TIDY ${CLANG_TIDY_COMMAND}
  )
endif()
```

# Motivation

# The downsides of using clang-tidy

- Running static analysis takes time, a lot of time.
  - Often more time than the actual compilation.
  - Unacceptable increase in build times.
  - Especially in CI pipelines doing full rebuilds.
  - Switching static analysis on/off:
    - Not ideal.
    - When do we flip the switch?
    - Typically ends up always off.



# Switching clang-tidy checks on/off with cmake

Add a cmake option:

```
option(  
  WITH_STATIC_ANALYSIS  
  "Enable static analysis" ON  
)
```

Add analysis-related target properties only when switched on:

```
if(WITH_STATIC_ANALYSIS and CLANG_TIDY_COMMAND)  
  set_target_properties(  
    my_target PROPERTIES  
    CXX_CLANG_TIDY ${CLANG_TIDY_COMMAND}  
  )  
endif()
```

## The reason for build slowdowns

- Often most of the code that is analysed doesn't change.
  - It is re-checked with the same result over and over.
  - Unless you change something in the “core” sources included everywhere.
  - Happens typically in CI pipelines, but occasionally also on developers' machines.

## The goal

- Have static analysis always on during development.
- Don't wait for rechecking of unchanged code.

# Solution

## The solution – analysis result caching

- If we can uniquely identify a static-analysis tool invocation we can store the result and retrieve it when the same invocation is repeated.
  - Similar to compilation-caching<sup>3</sup>.
  - Track all inputs of the analysis:
    - configuration options,
    - command-line arguments,
    - the source files,
    - etc.

---

<sup>3</sup><https://ccache.dev/>

◀ ◻ ▶ ◀ ◻ ▶ ◀ ≡ ▶ ◀ ≡ ▶ ≡ 🔍 ↻

## How does it work?

- Scans the inputs of `clang-tidy`:
  - command-line arguments,
  - configuration files,
  - analysed source files<sup>4</sup>.
- Makes a hash uniquely identifying the invocation from the above.

---

<sup>4</sup>preprocessed by the C/C++ preprocessor

## How does it work? (cont.)

- Checks if the hash is in the cache database:
  - if it is
    - doesn't run `clang-tidy` and returns immediately,
    - this is typically much faster.
  - otherwise
    - runs `clang-tidy` and if successful<sup>5</sup> store the hash.
    - This means that sources with findings keep being re-checked and the findings are shown.

---

<sup>5</sup>if there are no warnings or errors reported



## How is it used?

Create<sup>6</sup> a wrapper script, like:

```
#!/bin/bash  
REAL_CT=/full/path/to/clang-tidy  
  
/path/to/clang-tidy-cache \  
"${REAL_CT}" "${@}"
```

Put it into a directory in search `PATH`, before real `clang-tidy`:

```
export PATH="/path/to/wrapper-script-dir:${PATH}"
```

---

<sup>6</sup>or use the one in the repository

## Modes of operation

- Local
- Client / Server

## Local mode

- Stores the database in a local directory hierarchy.
- Location determined by the `CTCACHE_DIR` environment variable.
- By default a sub-tree in the temporary directory.
- If you want persistence, specify a directory in a disk-based file system.

## Client / server mode

- `clang-tidy-cache-server`
  - HTTP server exposing a REST API.
  - Can be used to store and retrieve hashes from the client.
  - The client (`clang-tide-cache`) can query the server – still way faster than running `clang-tidy`.

# Service – Rest API

- `http://ctcache:5000/...`
  - `/cache/<hash>` – insert `<hash>` into cache.
  - `/is_cached/<hash>` – tests if `<hash>` is cached.
  - `/purge_cache` – remove all cached hashes.
  - `/info` – static configuration information<sup>7</sup>.
  - `/stats` – server run-time status information<sup>8</sup>.
  - `/stats/*` – individual server status readouts<sup>9</sup>.
  - `/stats/ctcache.json` – long-term persistent server status information<sup>10</sup>.
  - `/images/*` – status chart images<sup>11</sup>.

---

<sup>7</sup>as JSON object

<sup>8</sup>as JSON object

<sup>9</sup>as JSON values

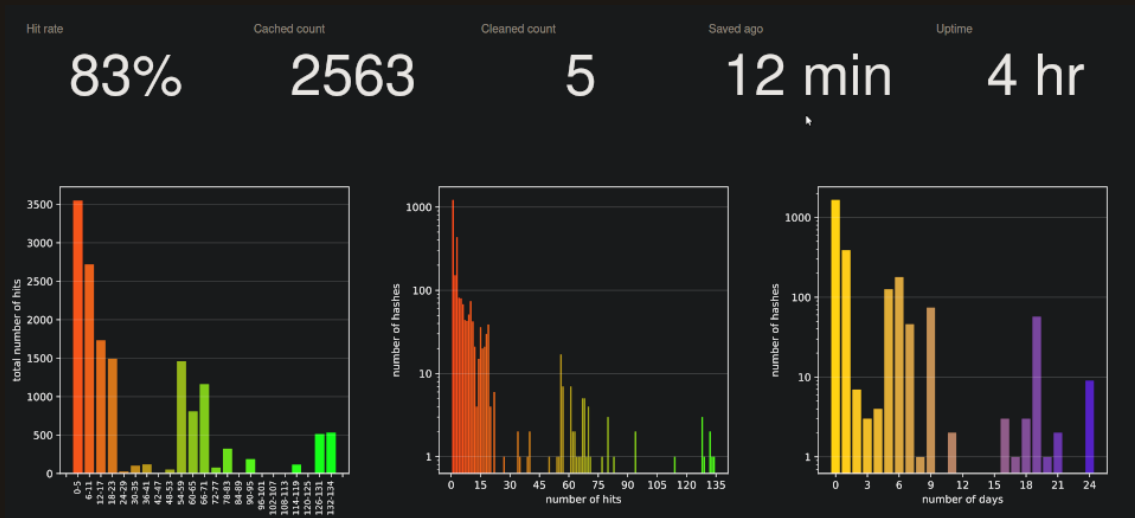
<sup>10</sup>as JSON file

<sup>11</sup>as SVG

## Server pages

- The clang tidy-cache's HTTP server also serves several web pages that are designed to be viewed in a browser:
  - the *dashboard* – the main one,
  - SVG plots showing various server statistics.

# Server dashboard



# Deployment



## Deploying the server

- There are several ways how to run the server:
  - just run it in Python,
  - as a *systemd* service,
  - in a *docker* container.

# Using python3

If you want to try it out:

```
python3 ./clang-tidy-cache-server
```

Check command-line arguments:

```
python3 ./clang-tidy-cache-server --help
```

```
usage: clang-tidy-cache-server
  [-h] [--debug] [--port NUMBER]
  [--save-path FILE-PATH.gz]
  [--save-interval NUMBER]
  [--stats-save-interval NUMBER]
  [--cleanup-interval NUMBER]
  [--stats-path DIR-PATH]
  [--chart-path DIR-PATH]
```

# Using systemd – installation

Install server as user service to home directory:

```
cd path/to/ctcache_repo  
./install-user-service
```

BTW: install client to user's home directory:

```
cd path/to/ctcache_repo  
./install-user-client
```

Installed files:

```
~/local/bin/clang-tidy # default wrapper script  
~/local/bin/clang-tidy-cache  
~/local/bin/clang-tidy-cache-server  
~/local/share/ctcache/static/* # static web files  
~/config/ctcache/systemd_env # systemd environment  
~/config/systemd/user/ctcache.service
```

# Using systemd – service start/stop

Reload the user service files:

```
systemctl --user daemon-reload
```

Start the service:

```
systemctl --user start ctcache.service
```

Stop the service:

```
systemctl --user stop ctcache.service
```

# Using systemd

Permanently enable automatic start of the service:

```
systemctl --user enable ctcache.service
```

Permanently disable automatic start of the service:

```
systemctl --user disable ctcache.service
```

# Using docker

## Build the image:

```
docker build -t ctcache .
```

## Basic usage:

```
docker run \  
  -e CTCACHE_PORT=5000 \  
  -p "80:5000" \  
  -it --rm \  
  --name ctcache ctcache
```

# Make cache data persistent with `docker volumes`

## Create the volume:

```
docker volume create ctcache
```

## Start container using the volume:

```
docker run \  
  -e CTCACHE_PORT=5000 \  
  -p "80:5000" \  
  -v "ctcache:/var/lib/ctcache" \  
  -it --rm \  
  --name ctcache ctcache
```

# Using docker-compose

The `docker-compose.yaml` file:

```
version: "3.6"
services:
  ctcache:
    build: .
    ports:
      - "5000:5000"
    volumes:
      - "ctcache:/var/lib/ctcache"
volumes:
  ctcache:
```



# Using docker-compose

Start the service container using docker-compose:

```
docker-compose up
```

as a daemon:

```
docker-compose up -d
```

Stop the running daemon and cleanup the container:

```
docker-compose down
```

# Environment variables

variable	client	server	meaning
CTCACHE_CLANG_TIDY	✓		path to the clang-tidy executable to be used
CTCACHE_DISABLE	✓		disables cache, always runs clang-tidy
CTCACHE_SKIP	✓		disables analysis, client returns "OK" immediately
CTCACHE_STRIP	✓		list of strings stripped from hashed inputs
CTCACHE_DIR	✓		the cache directory in local mode
CTCACHE_HOST	✓	✓	hostname or IP address of the server
CTCACHE_PORT	✓	✓	port number on which the server accepts connections
CTCACHE_WEBROOT		✓	directory where static served files are located

# Measurements

# Test projects

- Project 1 (small):
  - proprietary C++ training examples,
  - some use **heavy** template meta-programming,
  - part of the code is not analyzed,
  - $\approx 6300$  LOC, *tens* of build targets.
- Project 2 (medium):
  - open-source C++ wrapper for EGL, OpenGL, OpenAL, ...,
  - <https://github.com/matus-chochlik/oglpplu2>,
  - $\approx 135k$  LOC, *hundreds* of build targets.
- Project 3 (large):
  - proprietary, production code for embedded HW,
  - there are some unfixed findings – some sources are re-checked,
  - $\approx 750k$  LOC, *thousands* of build targets.

# Test hardware setup

- development laptop:
  - i5-7200U @ 2.50GHz (4 cores),
  - 16GB RAM,
  - 250GB SSD.
- ctcache server:
  - RPi 3B,
  - BCM2837 64bit @ 1.20GHz (4 cores),
  - 1GB RAM,
  - 1TB USB HDD.
- connected over WiFi (5GHz /  $\approx 650\text{Mbps}^{12}$ ).

---

<sup>12</sup>according to `iwconfig`

# Test configurations

- “compiler & clang-tidy”
  - `export CCACHE_DISABLE=1`
  - `export CTCACHE_DISABLE=1`
- “ccache & clang-tidy”
  - `unset CCACHE_DISABLE`
  - `export CTCACHE_DISABLE=1`
- “compiler & ctcache”
  - `export CCACHE_DISABLE=1`
  - `unset CTCACHE_DISABLE`
- “ccache & ctcache”
  - `unset CCACHE_DISABLE`
  - `unset CTCACHE_DISABLE`


# Test execution

- For each test project:
  - `cd /path/to/build/dir`
  - `rm -rf ./`
  - `cmake ... /path/to/project/source`
  - do initial build with caches enabled<sup>13</sup> (no measurements),
  - for each test configuration:
    - setup environment variables,
    - `make clean`
    - `time make -j N ...`

---

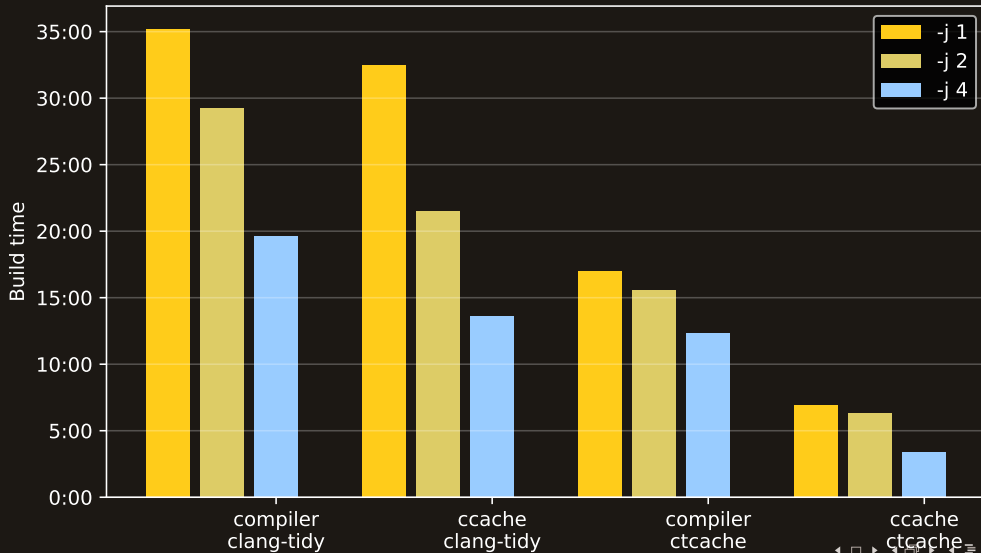
<sup>13</sup>to fill them

## Project 1 – clean build times (compilation caused some swapping on disk)

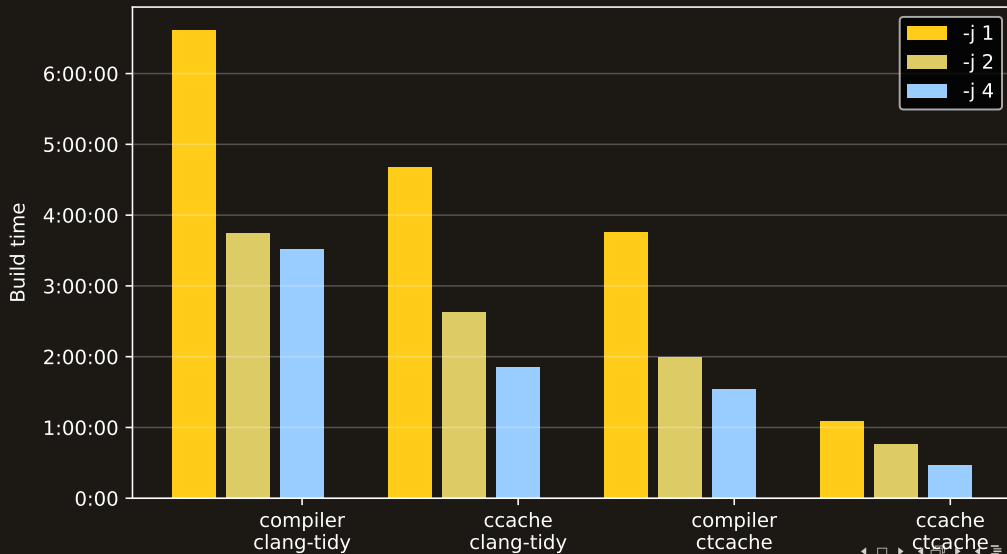





## Project 2 – clean build times




## Project 3 – clean build times



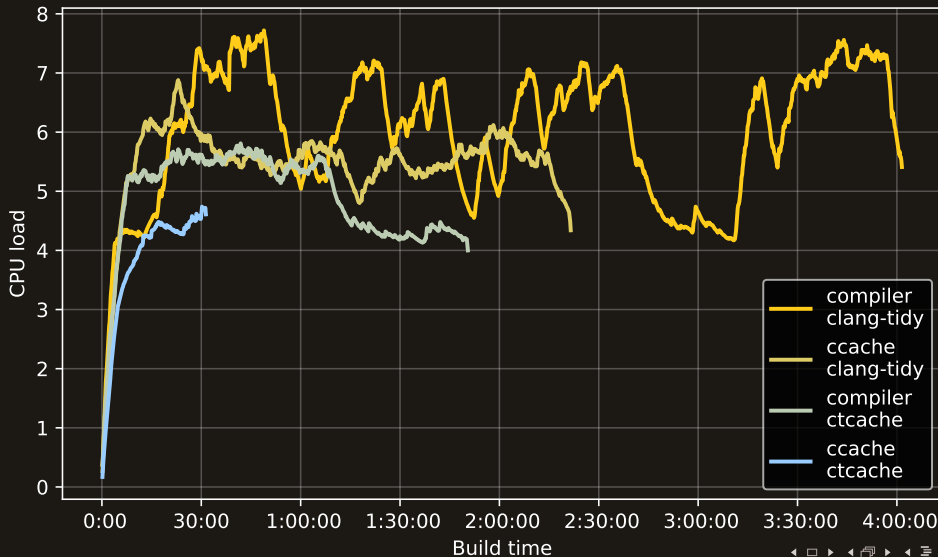
# Clean build speedups – $t_{uncached}/t_{cached}$



## Project 2 – system load (-j 4)



## Project 3 – system load (-j 4)



# Conclusions

# Conclusions

- Having static analysis (almost) always on is useful during development.
- Unnecessary delays from re-checking of unchanged code can be avoided by caching.
- Using `clang-tidy-cache` (and `ccache`) significantly improves build times<sup>14</sup> for projects of various sizes.
- The achieved speedups<sup>15</sup> are in the range of  $\approx 5x - 15x$ .

---

<sup>14</sup>especially clean rebuild times

<sup>15</sup>when using both `ccache` and `ctcache`

# Overview

- `clang-tidy-cache` (client and server)
  - are reasonably simple to setup,
  - provide many configuration options,
  - provide several deployment options,
  - can be integrated into CI pipelines<sup>16</sup>,
  - work well together with `ccache`,
  - can be shared among multiple users<sup>17</sup>.

---

<sup>16</sup>Jenkins, Travis, etc.

<sup>17</sup>more in a moment



## Personal experience

- Did development with `clang-tidy` checks enabled on “Project 2” and “Project 3” for more than a year.
- This helps to find and fix many bugs early.
- Also helps to enforce coding guidelines.
- With `clang-tidy-cache` the build times are kept to very reasonable levels.
- Build times increase only in rare cases when one of the “core” headers are changed.

## Possible future improvements

- Support other command-line static analysis tools.
- Support for *HTTPS* in the server.
- Number of requests over time statistics chart.
- Additional command-line options in the client:
  - Display server cache statistics on the command-line.
  - Reset server cache from the command-line.
- ...

# Thank you!

# Questions?

`https://github.com/matus-chochlik/ctcache`

`https://github.com/matus-chochlik/ctcache/doc/overview.pdf`

# Extras

## Sharing cache server among multiple users

- If multiple developers are building the same set of sources<sup>18</sup>,
- and they have reasonably similar development environments<sup>19</sup>,
- they have access to the same `clang-tidy-cache-server` instance,
- and they setup the `CTCACHE_STRIP` variable properly,
- they can share each others, cached static analysis results.

---

<sup>18</sup>they are working on the same project

<sup>19</sup>same platform, same compiler, STL and library versions

# What prevents analysis result sharing

- The cache works by hashing input strings – command-line arguments, pre-processed source file lines, configuration file lines, etc.
- The goal is to get the same hash for the “same” check.
- Between users with similar environments, the analysis inputs typically differ only in username-dependent sub-strings<sup>20</sup>.

---

<sup>20</sup>i.e. paths in the `-L`, `-I`, etc. compiler options

## The CTCACHE\_STRIP variable

- `CTCACHE_STRIP` – List of colon-separated strings, which are removed from the input strings.
  - For example  
`CTCACHE_STRIP="/home/user/myproject:/opt/custom/libs"`.
- If the “right” strings are stripped by every user → the same hashes are generated.
- May be somewhat tricky to setup properly.
- If set-up incorrectly may lead to false positives!